

AI Engineer Intern – Take-Home Assignment

Prompt Engineering & RAG Mini Project

Objective

Build a **small Retrieval-Augmented Generation (RAG) system** and demonstrate your ability to:

- Design **effective prompts**
 - Improve response quality through **iteration**
 - Evaluate and reason about LLM outputs
-

Problem Statement

You are given a set of **company policy documents** (e.g., Refund Policy, Cancellation Policy, Shipping Policy).

Your task is to build a **question-answering assistant** that:

1. Retrieves relevant information from the documents
 2. Generates **accurate, grounded answers**
 3. Avoids hallucinations
 4. Uses **clear and well-structured prompts**
-

Scope & Constraints

- Time limit: **4–6 hours**
- This is **not** a UI task (CLI or simple script is enough)
- Focus on **prompt quality, retrieval, and evaluation**
- You may use **OpenAI, Anthropic, or any open-source model**
- Use **Python**

Core Requirements

1. Data Preparation

- Load the provided policy documents (PDF / TXT / Markdown).
 - Clean and chunk the text using a reasonable strategy.
 - Explain **why** you chose your chunk size.
-

2. RAG Pipeline

Implement a basic RAG flow:

- Embedding generation
 - Vector storage (Chroma / Qdrant / FAISS)
 - Semantic retrieval (top-k)
 - Pass retrieved context to the LLM
-

3. Prompt Engineering (Very Important)

Design prompts that:

- Clearly instruct the model to answer **only from retrieved context**
- Handle **missing information gracefully**
- Use a **structured format** (headings, bullet points, citations, or JSON)

Include:

- Your **initial prompt**
 - At least **one improved iteration**
 - A short explanation of **what changed and why**
-

4. Evaluation

Create a small evaluation set (5–8 questions):

- Some answerable from the documents
- Some partially answerable
- Some unanswerable

Evaluate:

- Accuracy
- Hallucination avoidance
- Answer clarity

You can use:

- Manual evaluation
 - Simple scoring rubric ( /  / )
-

5. Edge Case Handling

Show how your system responds when:

- No relevant documents are found
 - The question is outside the knowledge base
-

Deliverables

1. GitHub repository containing:

- Source code
- `README.md`

2. `README` should include:

- Setup instructions
- Architecture overview
- Prompt(s) used

- Evaluation results
 - Key trade-offs and improvements you'd make with more time
-

Optional Bonus (Nice to Have)

(Do **not** over-optimize—these are bonus)

- Prompt templating with LangChain / LangGraph
 - Simple reranking step
 - Output schema validation (JSON)
 - Comparison between **two prompt versions**
 - Logging or tracing (very basic)
-

Evaluation Criteria (What We Look For)

Area	What We Evaluate
Prompting	Clarity, structure, hallucination control
RAG	Correct retrieval & grounding
Thinking	Trade-offs, explanations
Code	Readability, simplicity
Evaluation	Ability to judge model quality

What We Are *Not* Evaluating

- Fancy UI
- Large datasets
- Perfect accuracy
- Advanced ML math

We care about **reasoning, clarity, and iteration**.

Submission Instructions

- Share a **GitHub repo link**
- Include a short note explaining:
 - What you're most proud of
 - One thing you'd improve next